

Adventures in Supercomputing with R

Lecture 2: Hardware & Software Overview

George Ostrouchov

Oak Ridge National Laboratory and University of Tennessee

2022/01/19 (updated: 2022-02-23)

Summary of Last Lecure

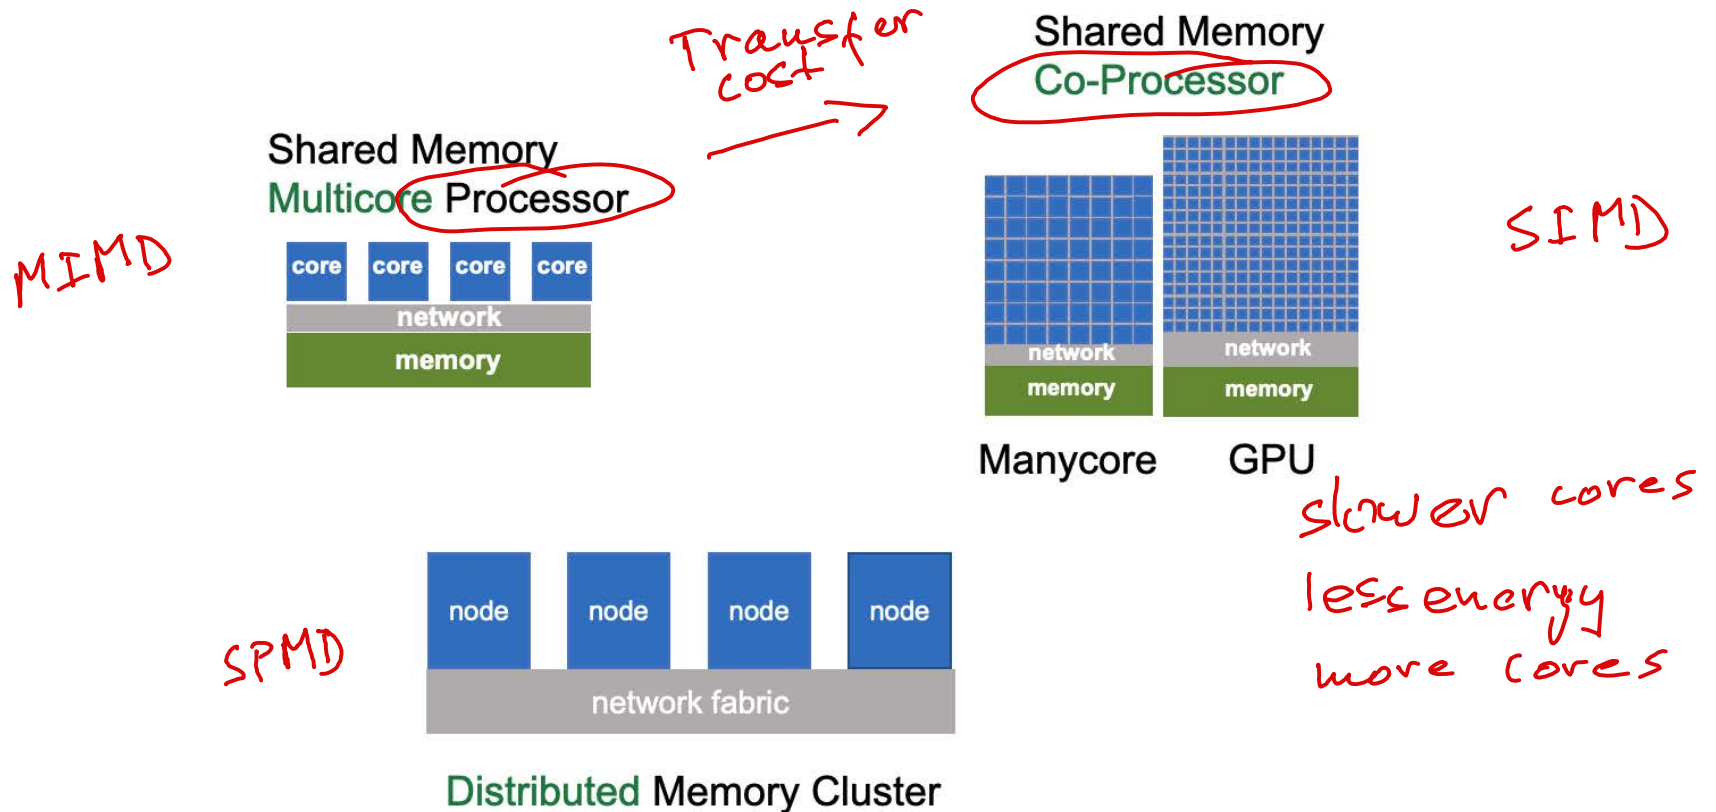
- Two parallel computing communities
 - Parallel numerical math (generating data) - supercomputing
 - Parallel databases (storing and serving data) - data centers
- Statistical computing left batch environment to be interactive
 - Developed S and then R
 - High-level, extensible, and interactive
 - Use numerical libraries
- Supercomputing takes R back to batch
 - High-level, extensible, but batch
 - Use scalable numerical libraries
- Workflow, accounts, tools
 - RStudio, git, GitHub, ssh, unix

PUTTY
Win SCP



Parallel Hardware

A high level look at what matters ...

Three Basic Concepts in Hardware




Three Basic Concepts in Hardware




Manycore

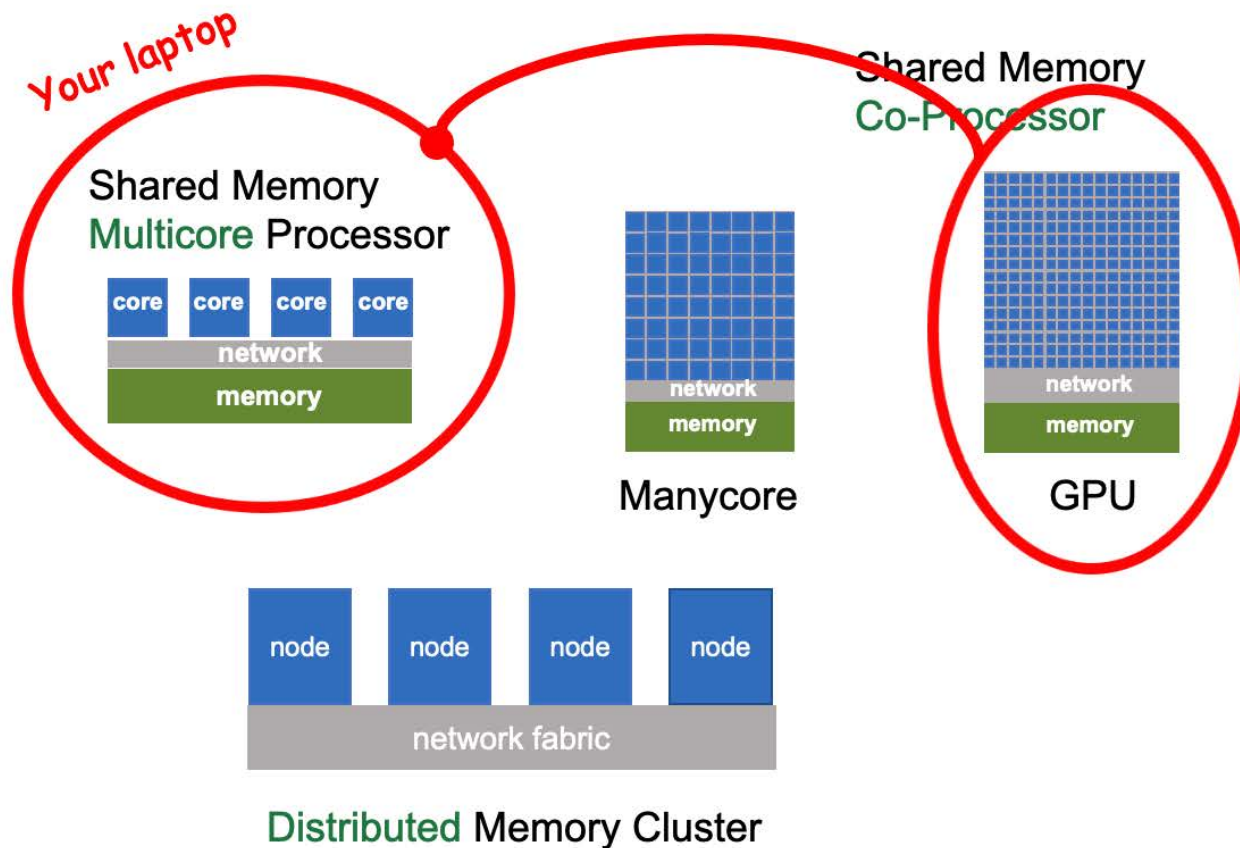
Shared Memory
Co-Processor



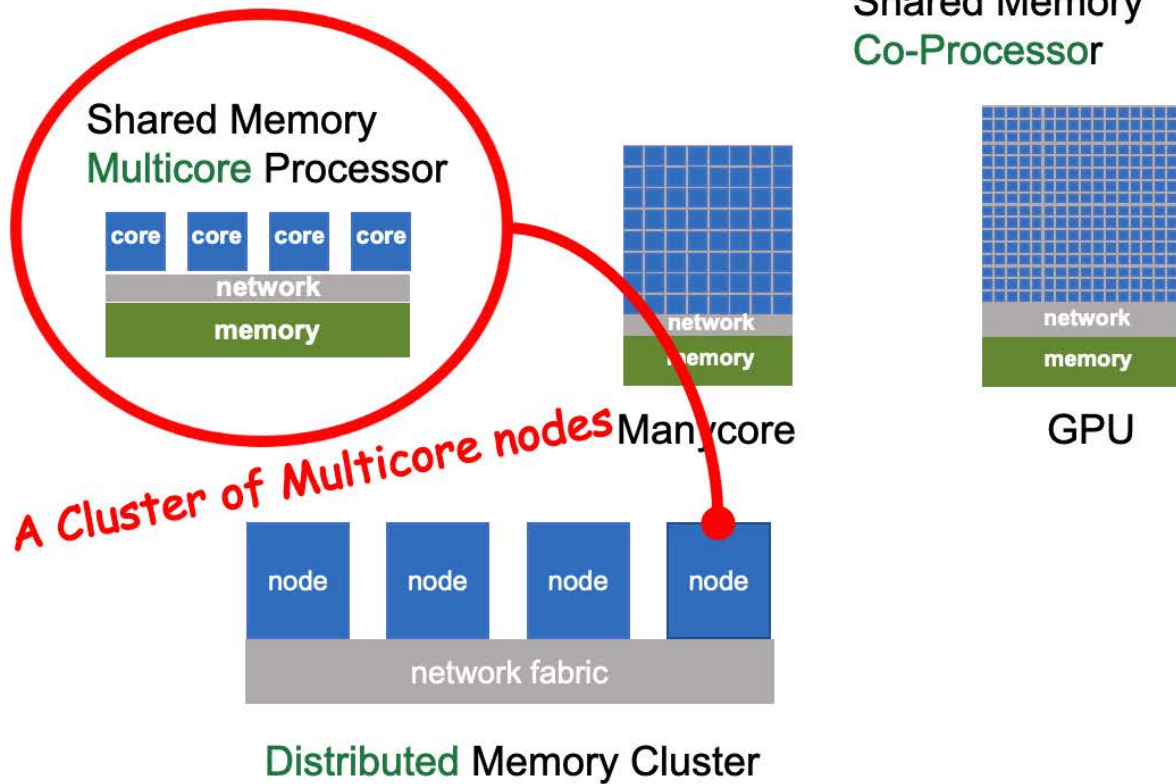
GPU

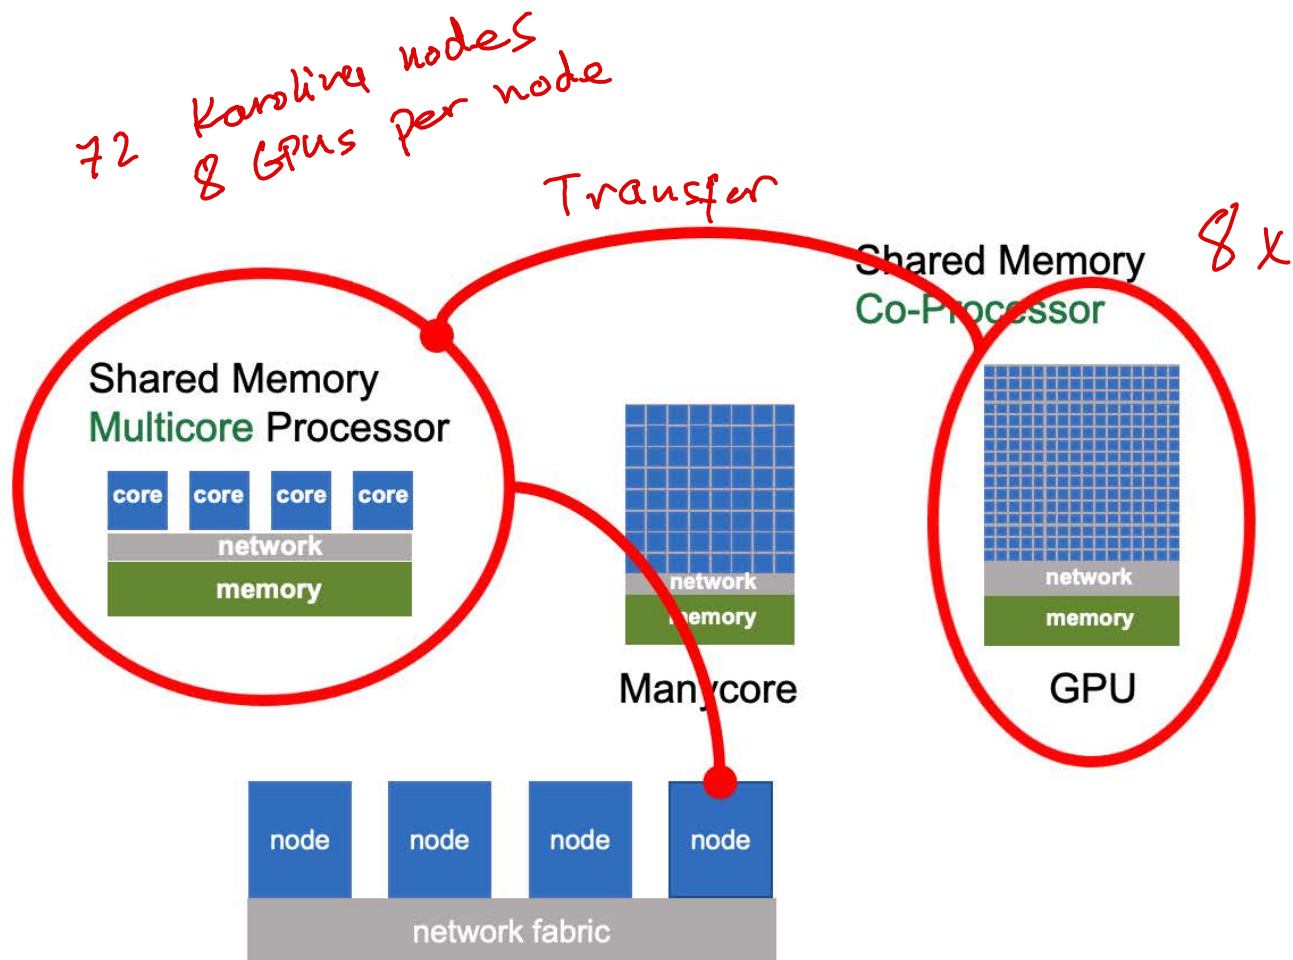


Distributed Memory Cluster

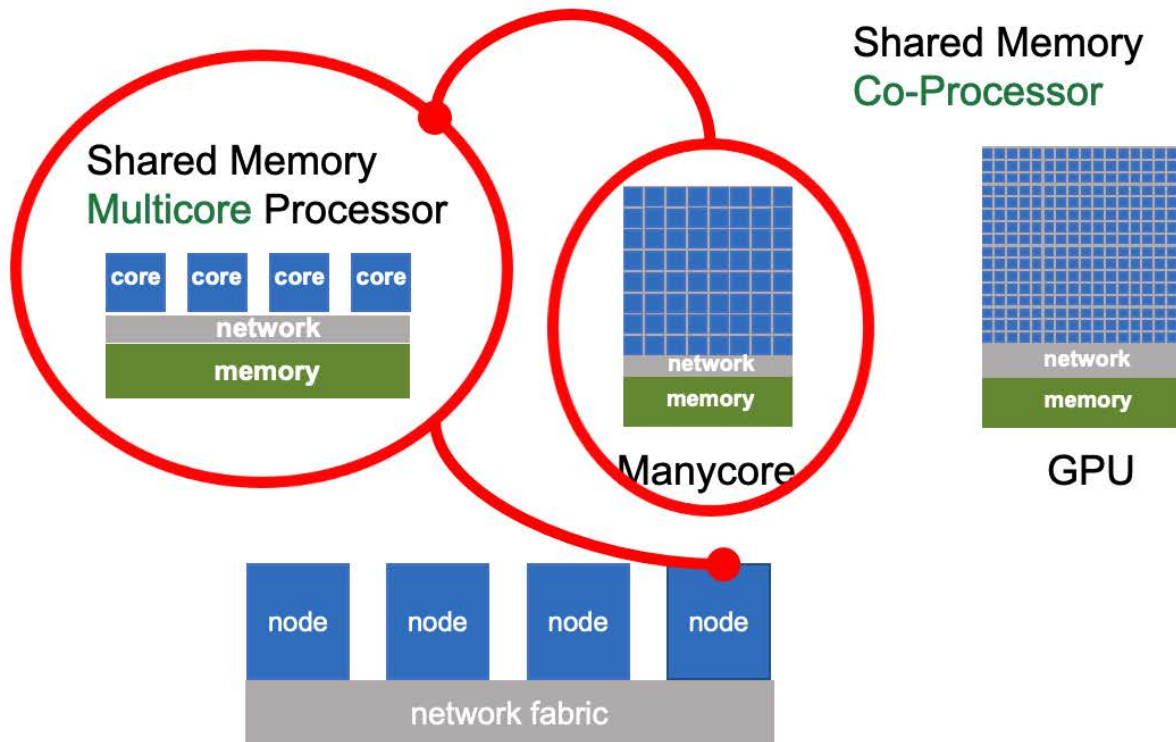


720 Karolina nodes
128 cores / node





Distributed Memory Cluster
A Cluster of Multicore nodes with GPU co-Processors



Distributed Memory Cluster
A Cluster of Multicore nodes with Manycore co-Processors




Best of both worlds. Future?

Can act like
↓

Programmable
like
↙


Shared Memory
Co-Processor

Shared Memory
Multicore Processor



GPU

Manycore



Distributed Memory Cluster

A Cluster of Manycore Nodes


Cray XC40
2nd generation Manycore
Intel Xeon Phi "KNL"

Fugaku ARM #1

Native Programming Mindset


Default is serial: which tasks should be made parallel?

Shared Memory
Multicore Processor




Offload data and tasks: We are slow but many!

Shared Memory
Co-Processor



GPU

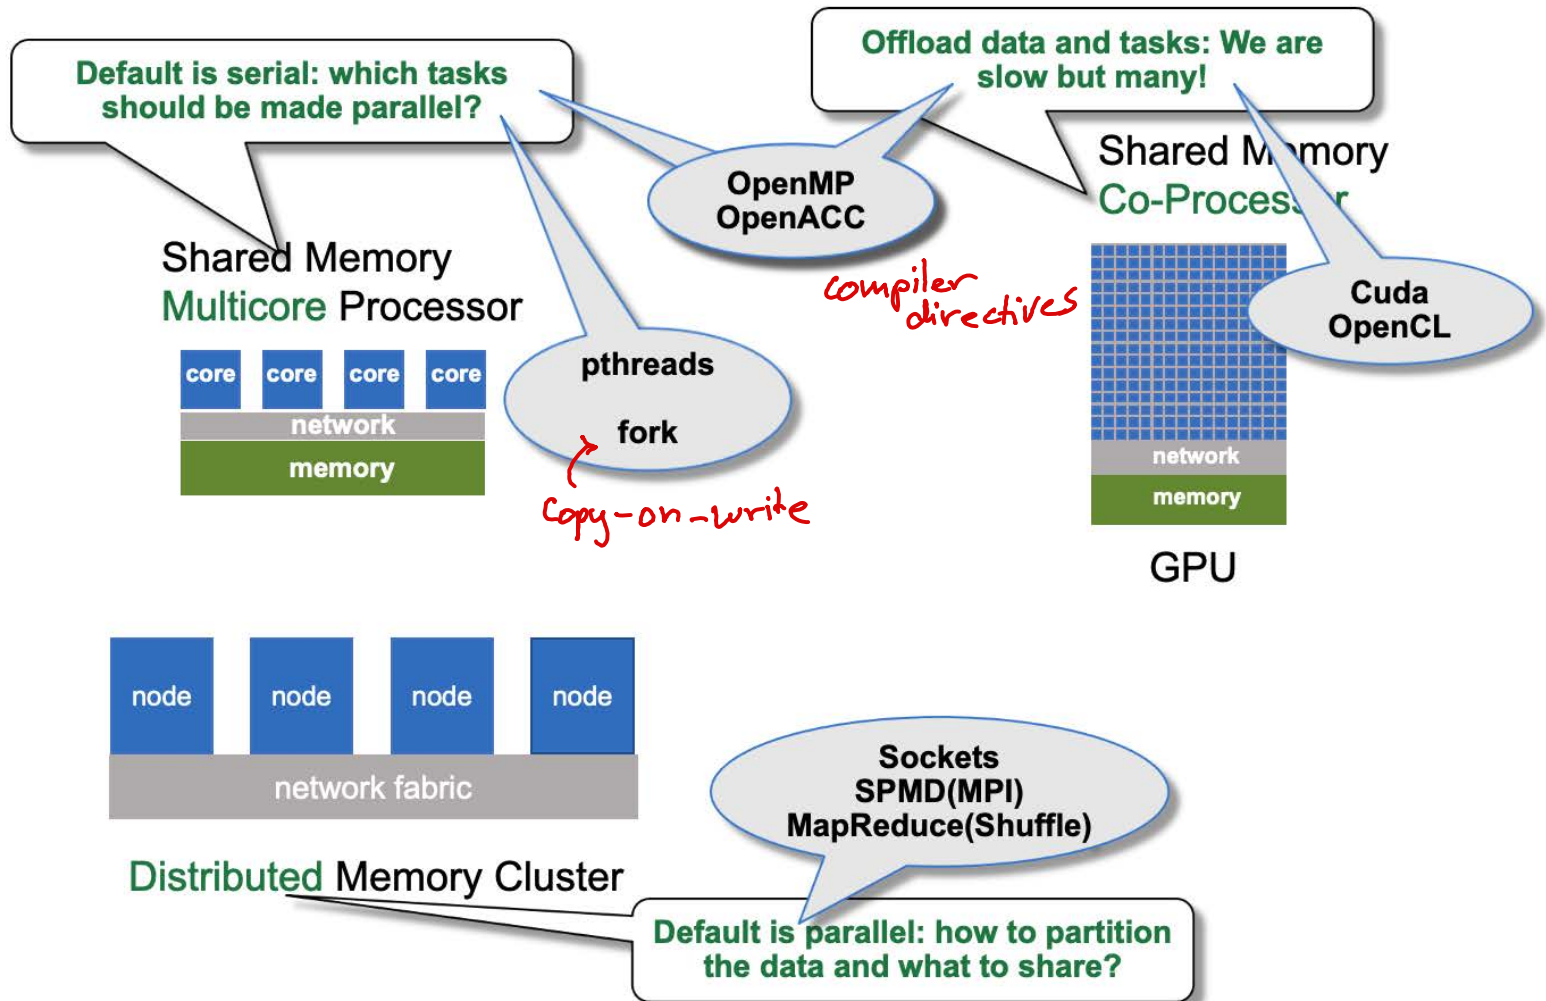


Distributed Memory Cluster

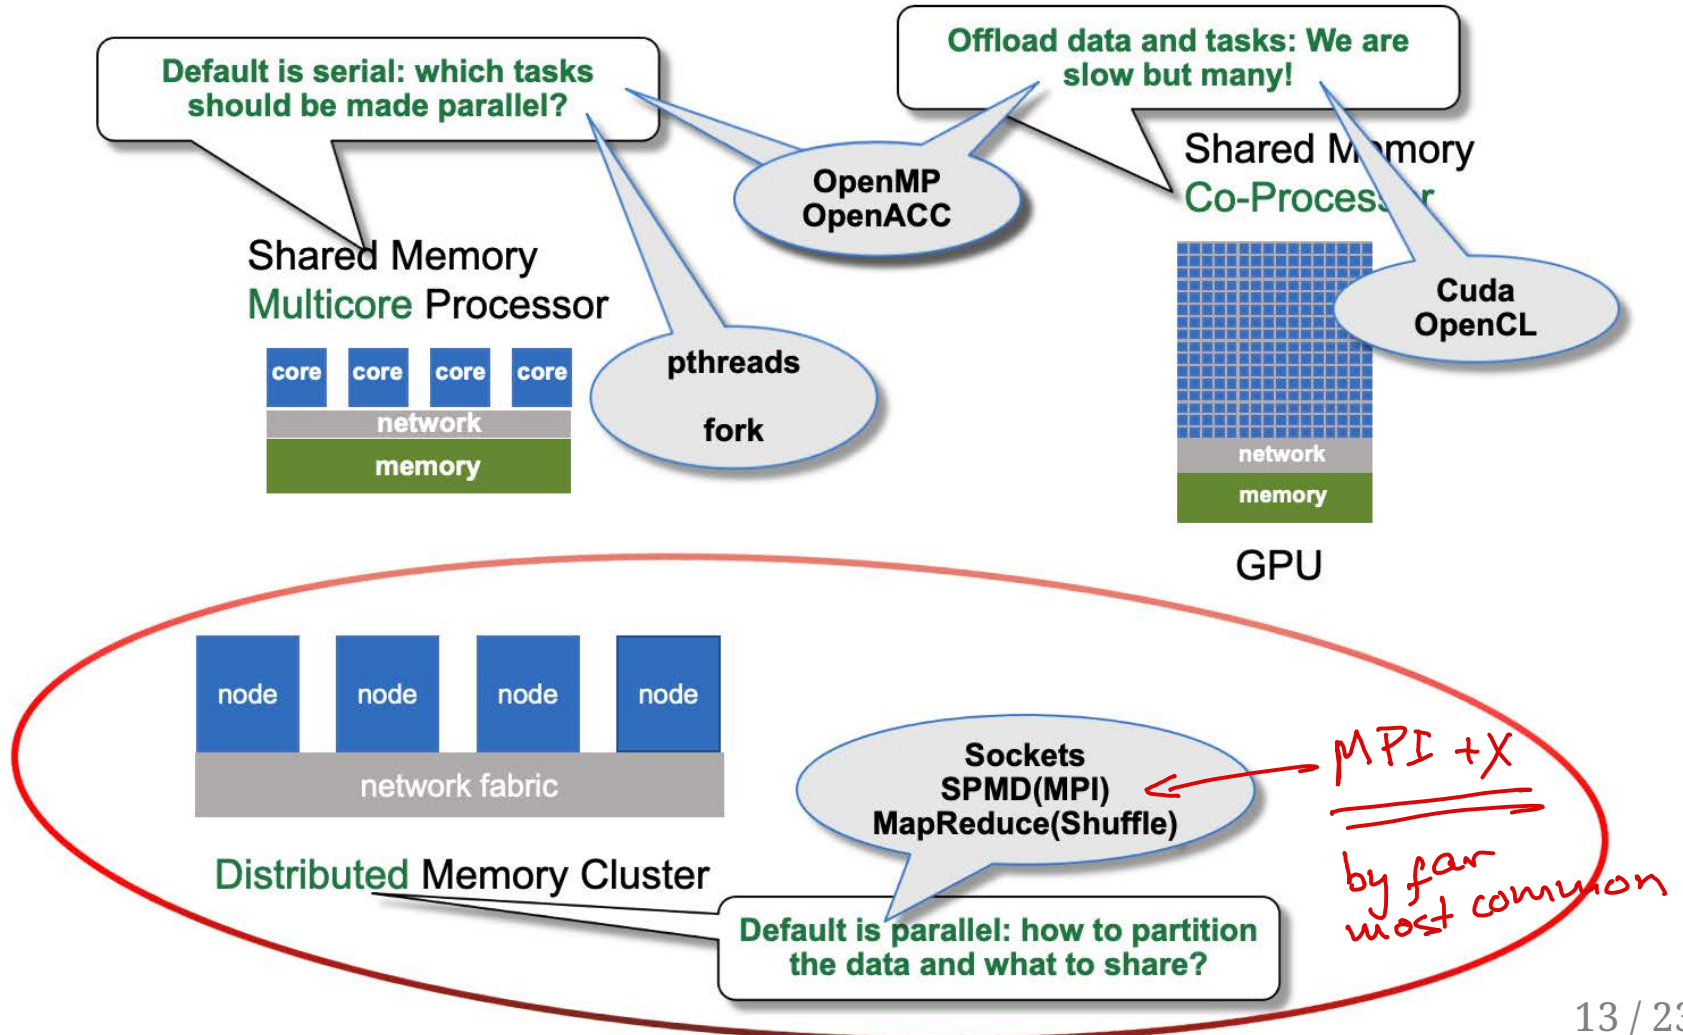
SPMD

Default is parallel: how to partition the data and what to share?

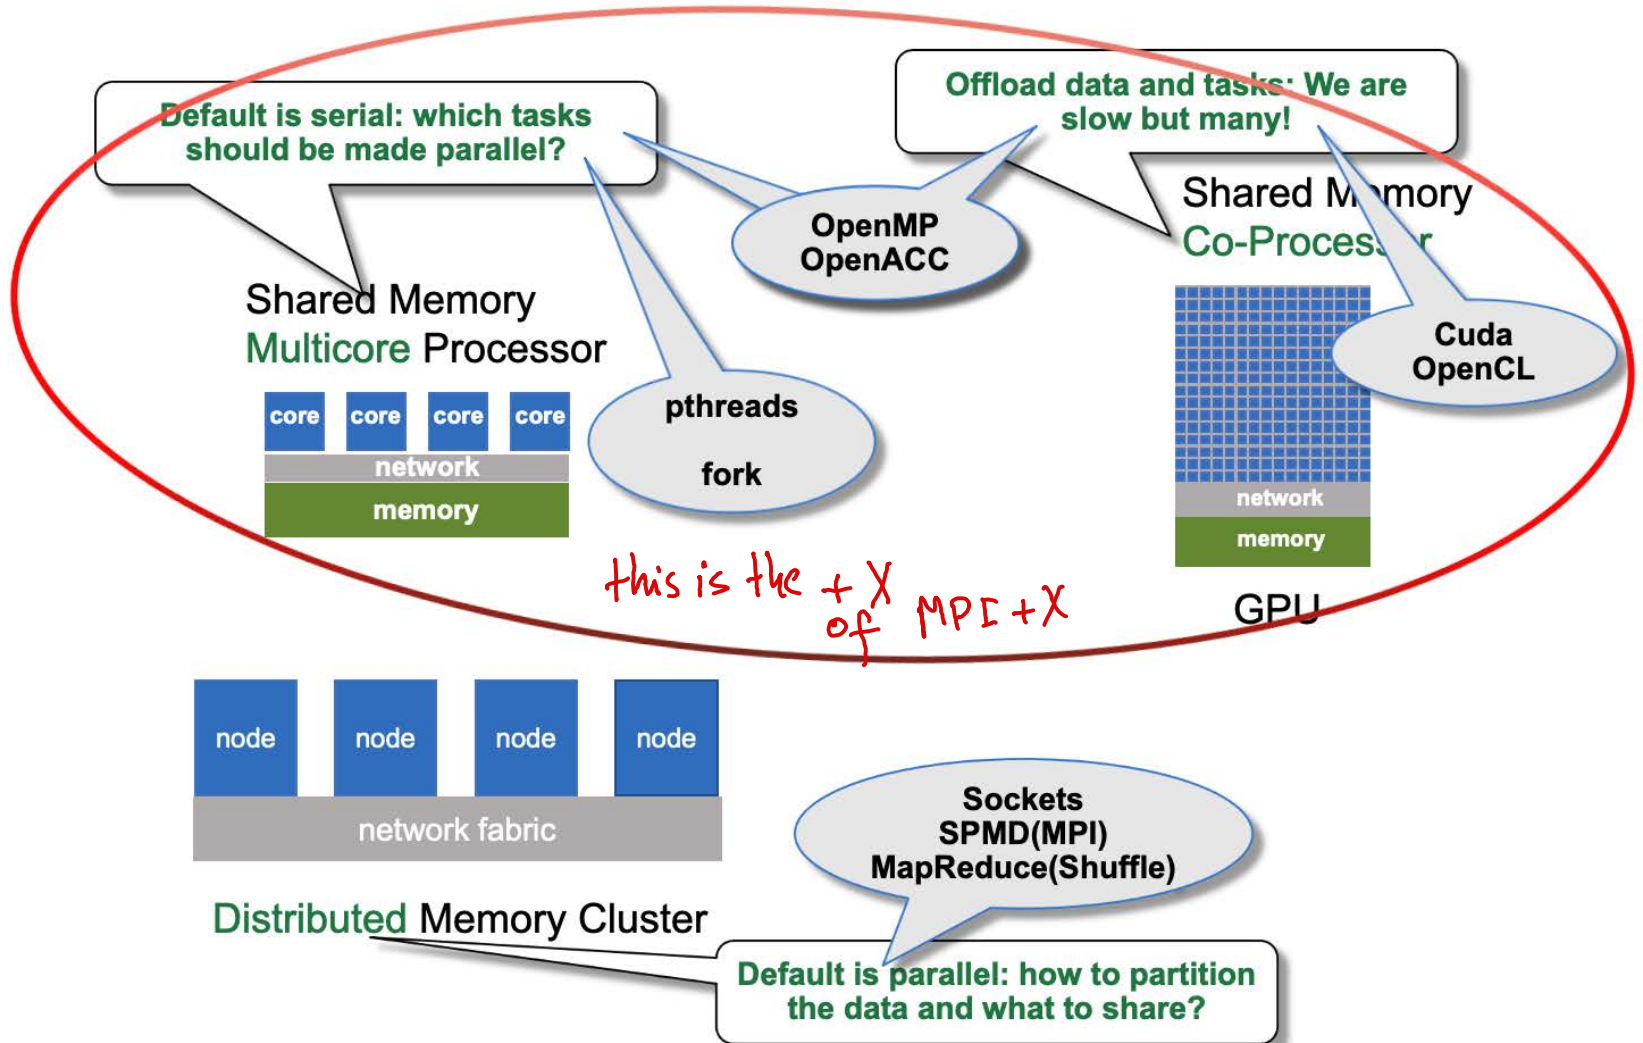
Native Programming Models and Tools



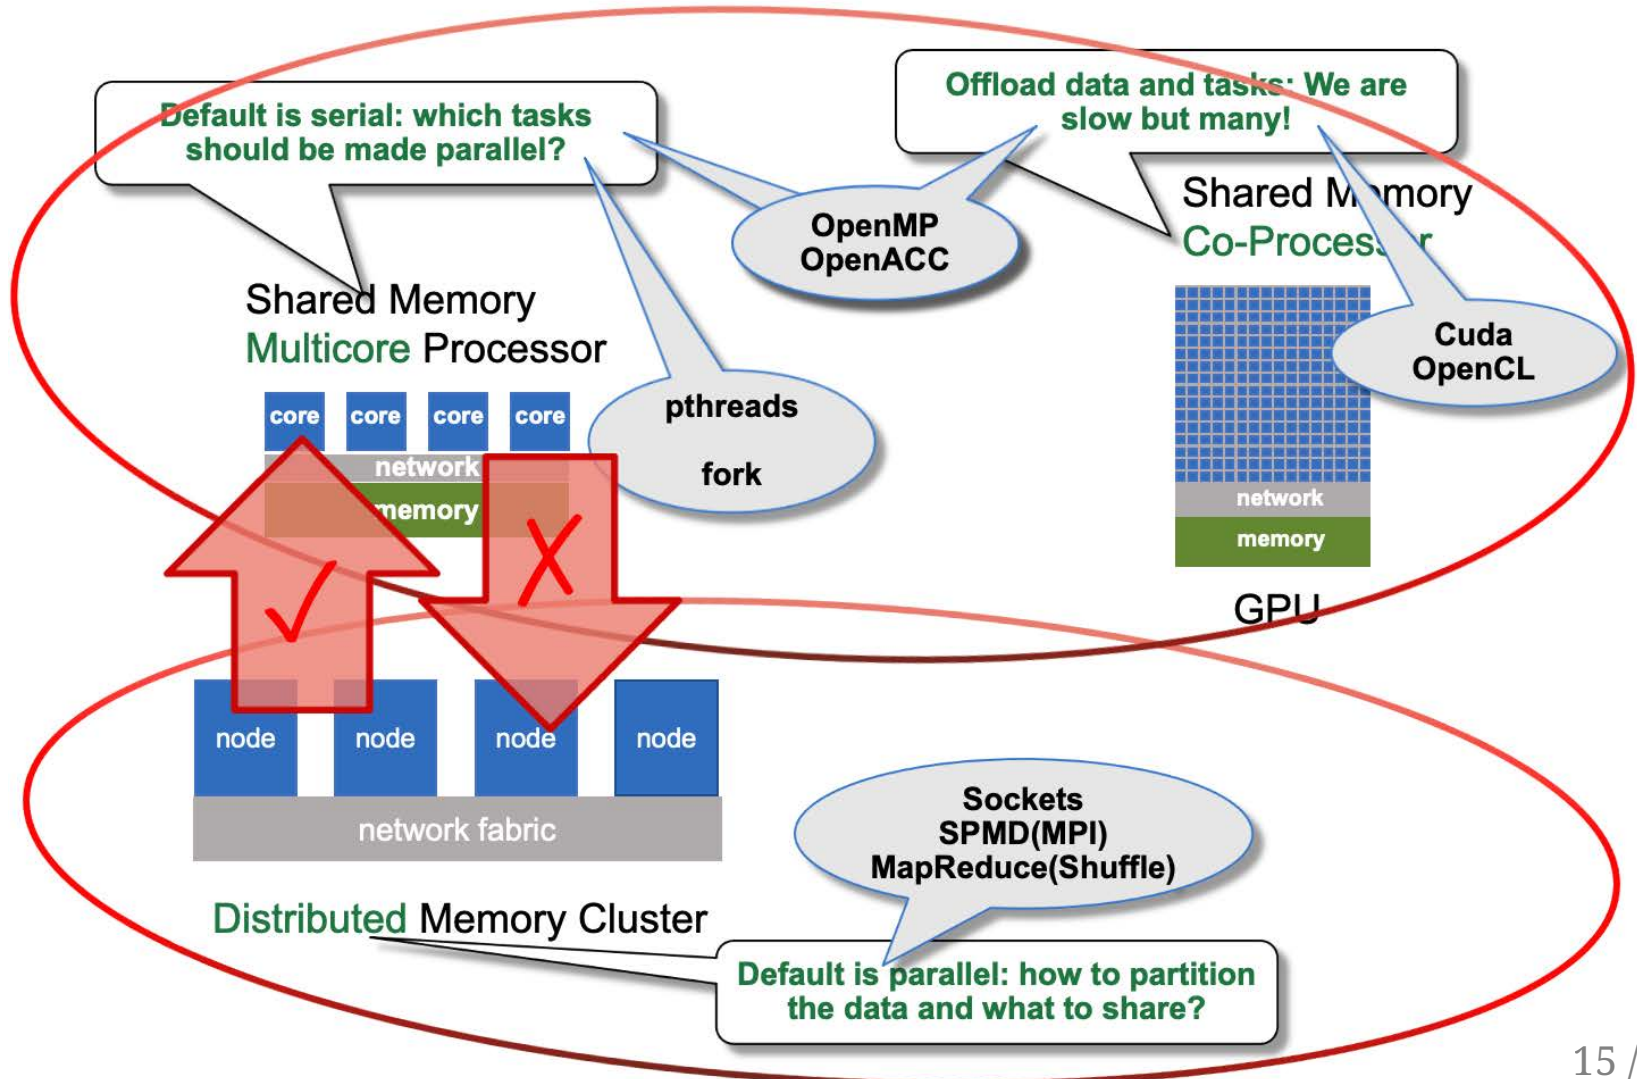
35+ Years of Parallel Computing Research



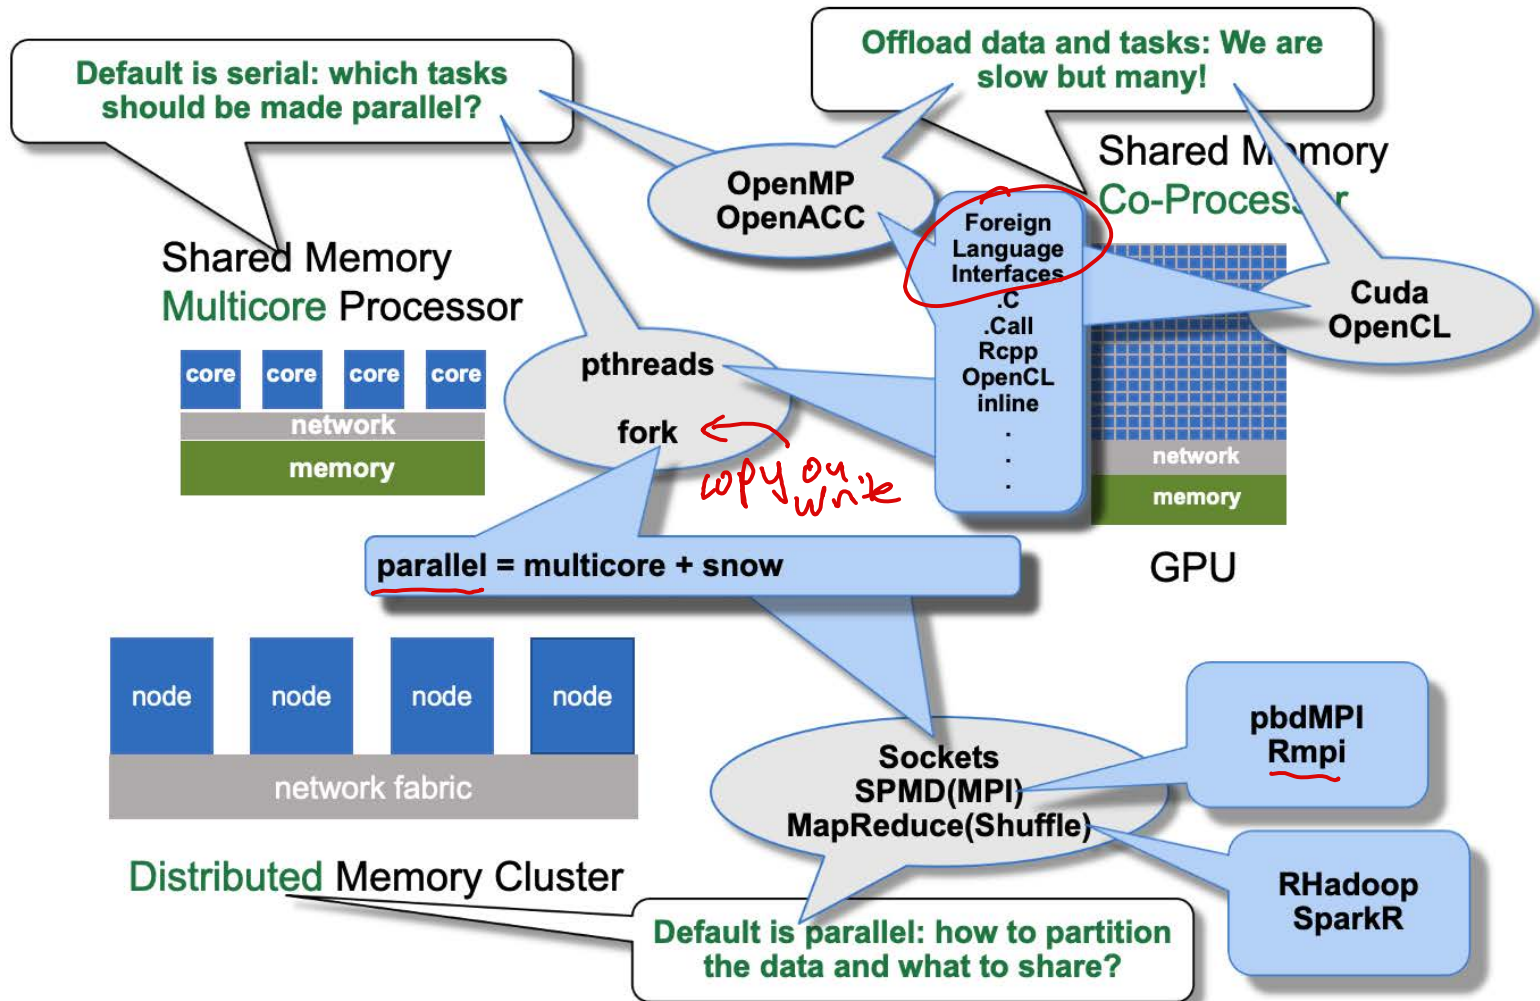
Last 15+ years of Advances



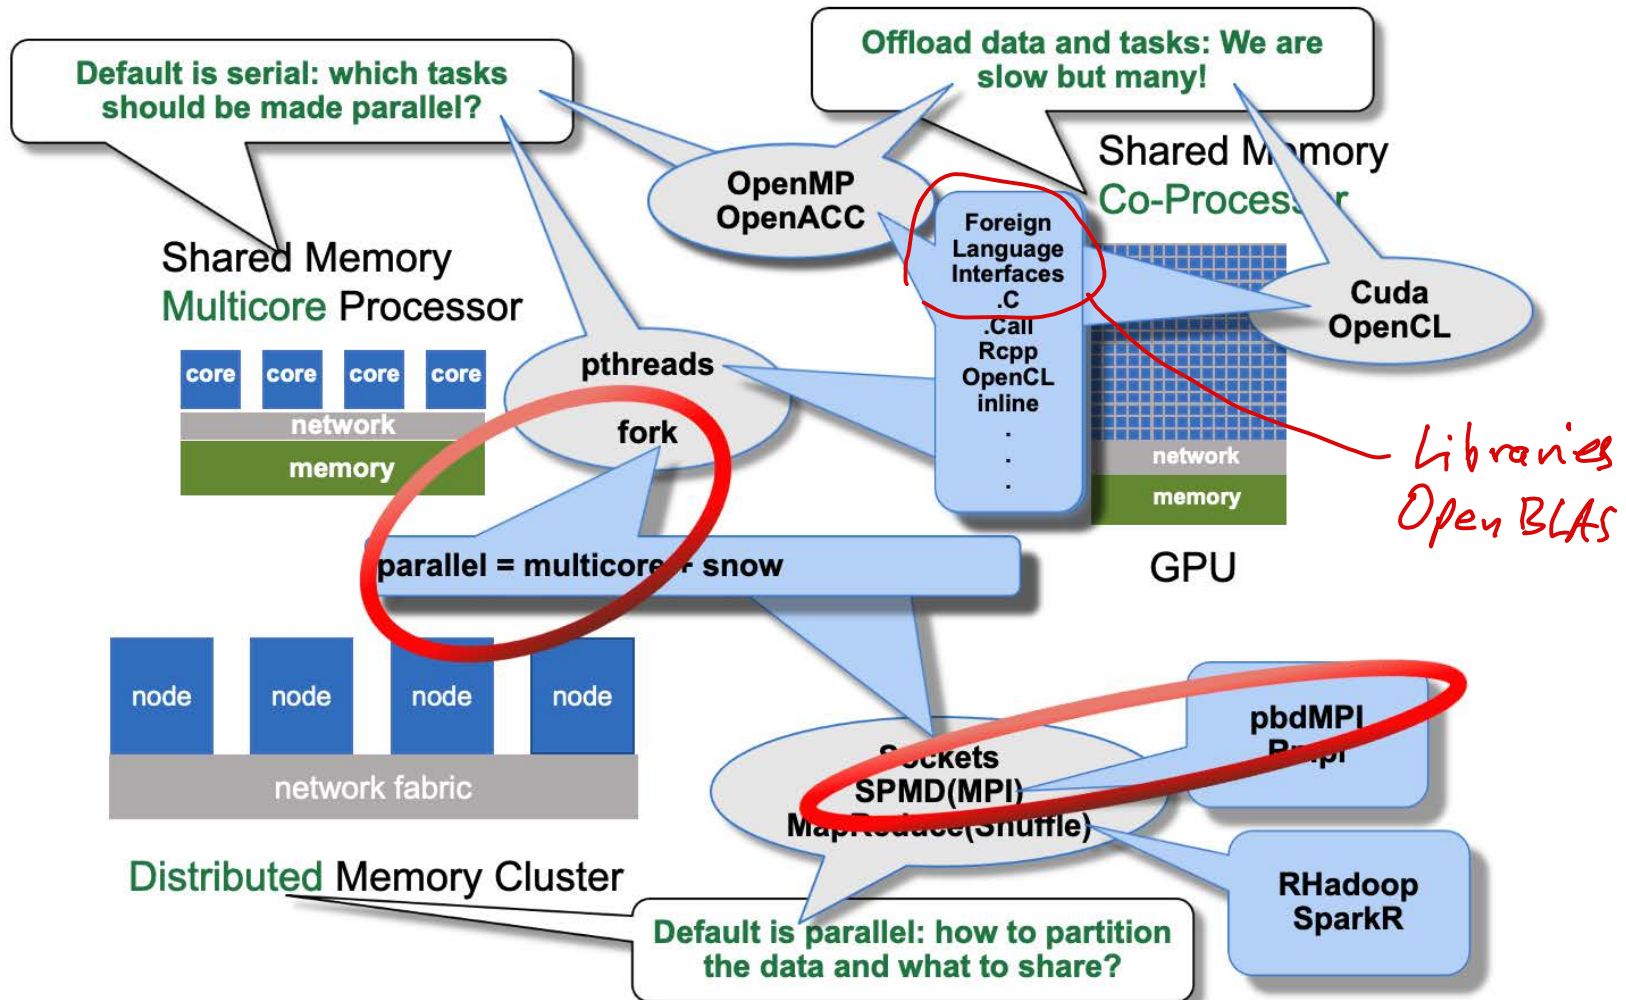
Distributed Programming Works in Shared Memory




R Interfaces to Low-Level Native Tools




R Interfaces to Low-Level Native Tools



Running Distributed on a Cluster



Typical Workflow from Laptop to Cluster



Unix Concepts and Starting your Workflow

...

Some useful Unix concepts

- Linux is one of many descendants of original Unix. Mac OS X is another.
- Like all file systems, Unix files are organized as a tree.
- When in a terminal, you are talking to a *shell* program (*bash* is most common)
- Commands are looked up in directories listed in your PATH variable
- Spaces delimit commands and options
- > and >> redirect standard output to a file
- *command1* | *command2* pipes standard output1 to standard input2
- \$ means substitute variable value *echo \$PATH*
- There are many resources on the web to learn Linux basics.

Some useful Linux commands

- **pwd** Show current directory
- **ls** List files in current directory
 - **ls -a** Include files that start with .
 - **ls -l** Long listing with *permissions*, *owners*, and *last change time*
- **cd** *dir_name* Change directory to *dir_name*
 - **cd** without *dir_name* or `~` as *dir_name* changes to home directory
- **mkdir** *dir_name* Creates directory *dir_name*
- **rmdir** *dir_name* Deletes directory (must be empty)
- **rm** *file_name* Deletes *file_name*
- **cat** *file_name* Displays content of entire *file_name*
- **less** *file_name* Displays content of *file_name* with paging
- • **man** *command* Displays the manual page for *command* with paging
- **which** *command* Returns location of *command*
- **exit** Quit shell and logout

Demo follows

Slides created via the R package **xaringan** and converted to pdf with **xaringanBuilder**