

Adventures in Supercomputing with R

Lecture 3: Parallel Speedup, Shared Memory Multicore

George Ostrouchov

Oak Ridge National Laboratory and University of Tennessee


2022/02/20 (updated: 2022-03-02)

Review of Last Lecture

- Shared memory and distributed memory
 - One program in both cases
 - Distributed programming works in shared memory
- Shared memory devices range between MIMD and SIMD extremes
 - Processor and co-processor
 - Multicore processors are MIMD and more easily programmable
 - GPUs co-processors are very efficient at SIMD
 - Manycore processors can be best of both worlds
- "Fork a repository" versus "Unix fork"
- There are other ways to edit remote code besides RStudio and GitHub


Measurement and terminology of parallel speedup ...

Embarrassingly (Pleasingly) Parallel



- t : Serial time
- n Number of chunks (or processes)
- t_n : Single chunk time with n chunks
- t_p : Parallel time


Serial Section (Amdahl's Law)



- t : Serial time (**fixed**)
- n Number of chunks (or processes)
- t_n : single chunk time with n chunks
- t_p : Parallel time
- t_s : Serial section time

Strong Scaling: fixed work, increasing resources


Serial Section (Gustafson's Law)



- t : Serial time (**growing**: $t_{2n} = 2t_n$)
- n : Number of chunks (or processes)
- t_n : single chunk time with n chunks
- t_p : Parallel time
- t_s : Serial section time


Weak Scaling: increasing work, increasing resources

Parallel Overhead



- t : Serial time
- n Number of processes
- t_n : single chunk time with n chunks
- t_p : Parallel time
- t_s : Serial section time
- $t_o(n)$: Parallel overhead time

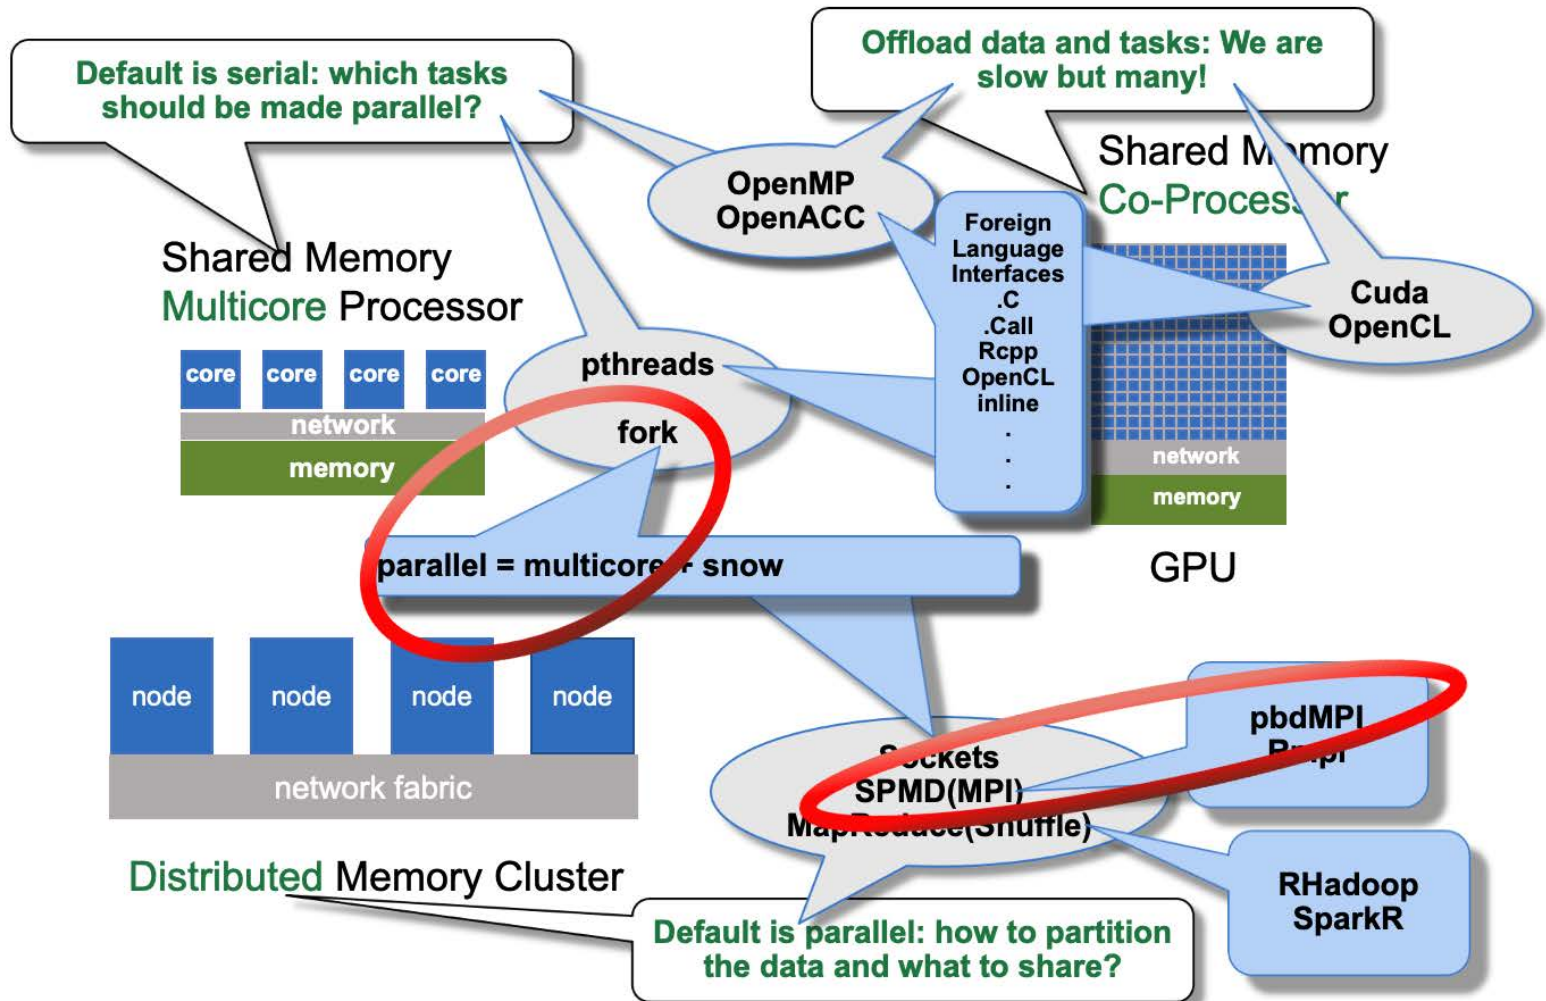
Preview - Distributed Compute Replication



- t : Serial time
- n Number of processes
- t_n : single chunk time with n chunks
- t_p : Parallel time
- t_s : Serial section time
- $t_o(n)$: Parallel overhead time

Replication can reduce communication overhead

R Interfaces to Low-Level Native Tools




Unix fork


- Copy-on-write
- A memory efficient parallelism on shared memory devices
- **parallel** package `mclapply` and friends
- Use for numerical sections only
 - Avoid GUI, I/O, and graphics sections
- Convenient for data (not modified)
- Convenient for functional languages like R
- Avoid or manage nested parallelism
 - OpenBLAS takes all cores by default
 - `data.table` automatically switches to single threaded mode upon fork

A deeper discussion of fork memory (if you have interest) on [YouTube](#) by Chris Kanich (UIC)

Copy-on-write



Shared memory after forking Process B



After Process B writes to copy-on-write shared memory page

Mapping Threads to Cores

Theory and Reality

- Operating system manages core affinity
- Operating system tasks can compete
- Core switching occurs frequently
- **But it works rather well!**



Drop-in replacements (almost) for `lapply`, `mapply`, and `Map`

```
mclapply(X, FUN, ..., mc.preschedule = TRUE, mc.set.seed =  
TRUE, mc.silent = FALSE, mc.cores = getOption("mc.cores", 2L),  
mc.cleanup = TRUE, mc.allow.recursive = TRUE, affinity.list =  
NULL)
```

```
mcmapply(FUN, ..., MoreArgs = NULL, SIMPLIFY = TRUE, USE.NAMES  
= TRUE, mc.preschedule = TRUE, mc.set.seed = TRUE, mc.silent =  
FALSE, mc.cores = getOption("mc.cores", 2L), mc.cleanup =  
TRUE, affinity.list = NULL)
```

```
mcMap(f, ...)
```

Example Random forest Code

Letter recognition data (20 000 \times 17)



```
[,1] lettr capital letter
[,2] x.box horizontal position of box
[,3] y.box vertical position of box
[,4] width width of box
[,5] high height of box
[,6] onpix total number of on pixels
[,7] x.bar mean x of on pixels in box
[,8] y.bar mean y of on pixels in box
[,9] x2bar mean x variance
[,10] y2bar mean y variance
[,11] xybar mean x y correlation
[,12] x2ybr mean of x^2 y
[,13] xy2br mean of x y^2
[,14] x.ege mean edge count left to right
[,15] xegvy correlation of x.ege with y
[,16] y.ege mean edge count bottom to top
[,17] yegvx correlation of y.ege with x
```

Figure 1: Letter Recognition data (image: [Frey and Slate, 1991](#)), description: **mlbench** package).

*Parallel Statistical Computing with R: An Illustration on Two Architectures
[arXiv:1709.01195](#)


Random Forest Classification

Build many decision trees from random subsets of variables

Use their majority votes to classify

Example Random Forest Classification Code

Letter recognition data (20 000 \times 17)



KPMS-IT4I-EX/code/rf_serial.r

```
library(randomForest)
data(LetterRecognition, package = "mlbench")
set.seed(seed = 123)

n = nrow(LetterRecognition)
n_test = floor(0.2 * n)
i_test = sample.int(n, n_test)
train = LetterRecognition[-i_test, ]
test = LetterRecognition[i_test, ]

rf.all = randomForest(lettr ~ ., train, ntree = 500, norm.votes = FALSE)
pred = predict(rf.all, test)

correct = sum(pred == test$lettr)
cat("Proportion Correct:", correct/(n_test), "\n")
```

KPMS-IT4I-EX/code/rf_mc.r

```
library(parallel)
library(randomForest)
data(LetterRecognition, package = "mlbench")
set.seed(seed = 123, "L'Ecuyer-CMRG")

n = nrow(LetterRecognition)
n_test = floor(0.2 * n)
i_test = sample.int(n, n_test)
train = LetterRecognition[-i_test, ]
test = LetterRecognition[i_test, ]

nc = as.numeric(commandArgs(TRUE)[1])
ntree = lapply(splitIndices(500, nc), length)
rf = function(x) randomForest(lettr ~ ., train, ntree=x, norm.votes =
rf.out = mclapply(ntree, rf, mc.cores = nc)
rf.all = do.call(combine, rf.out)

crows = splitIndices(nrow(test), nc)
rfp = function(x) as.vector(predict(rf.all, test[x, ]))
cpred = mclapply(crows, rfp, mc.cores = nc)
pred = do.call(c, cpred)


correct <- sum(pred == test$lettr)
cat("Proportion Correct:", correct/(n_test), "\n")
```

Pseudo Random Number Generators (RNG)


- Guaranteed reproducibility
- Possibly overlapping streams

- Reproducibility for same number of streams
- Guaranteed independent streams

Random seeds with serial RNG



Parallel RNG



KPMS-IT4I-EX/code/rf_mc.r

```
library(parallel)
library(randomForest)
data(LetterRecognition, package = "mlbench")
set.seed(seed = 123, "L'Ecuyer-CMRG")

n = nrow(LetterRecognition)
n_test = floor(0.2 * n)
i_test = sample.int(n, n_test)
train = LetterRecognition[-i_test, ]
test = LetterRecognition[i_test, ]

nc = as.numeric(commandArgs(TRUE)[1])
ntree = lapply(splitIndices(500, nc), length)
rf = function(x) randomForest(lettr ~ ., train, ntree=x, norm.votes =
rf.out = mclapply(ntree, rf, mc.cores = nc)
rf.all = do.call(combine, rf.out)

crows = splitIndices(nrow(test), nc)
rfp = function(x) as.vector(predict(rf.all, test[x, ]))
cpred = mclapply(crows, rfp, mc.cores = nc)
pred = do.call(c, cpred)

correct <- sum(pred == test$lettr)
cat("Proportion Correct:", correct/(n_test), "\n")
```

KPMS-IT4I- EX/code/rf_karolina_pbs.sh

```
#!/bin/bash
#PBS -N rf
#PBS -l select=1:ncpus=128,walltime=00:05:00
#PBS -q qexp
#PBS -e rf.e
#PBS -o rf.o

cd ~/KPMS-IT4I-EX/code
pwd

module load R
echo "loaded R"

time Rscript rf_serial.r
time Rscript rf_mc.r 1
time Rscript rf_mc.r 2
time Rscript rf_mc.r 4
```

Demo ...