

# Adventures in Supercomputing with R

## Lecture 3: Parallel Speedup, Shared Memory Multicore

George Ostrouchov

Oak Ridge National Laboratory and University of Tennessee


2022/02/20 (updated: 2022-03-02)

# Review of Last Lecture

- Shared memory and distributed memory
  - One program in both cases
  - Distributed programming works in shared memory
- Shared memory devices range between MIMD and SIMD extremes
  - Processor and co-processor
  - Multicore processors are MIMD and more easily programmable
  - GPUs co-processors are very efficient at SIMD
  - Manycore processors can be best of both worlds
- "Fork a repository" versus "Unix fork"
- There are other ways to edit remote code besides RStudio and GitHub


# Measurement and terminology of parallel speedup ...

# Embarrassingly (Pleasingly) Parallel



- $t$ : Serial time
- $n$  Number of chunks (or processes)
- $t_n$ : Single chunk time with  $n$  chunks
- $t_p$ : Parallel time


# Serial Section (Amdahl's Law)



- $t$ : Serial time (**fixed**)
- $n$  Number of chunks (or processes)
- $t_n$ : single chunk time with  $n$  chunks
- $t_p$ : Parallel time
- $t_s$ : Serial section time

**Strong Scaling: fixed work, increasing resources**


# Serial Section (Gustafson's Law)



- $t$ : Serial time (**growing**:  $t_{2n} = 2t_n$ )
- $n$ : Number of chunks (or processes)
- $t_n$ : single chunk time with  $n$  chunks
- $t_p$ : Parallel time
- $t_s$ : Serial section time


**Weak Scaling: increasing work, increasing resources**

# Parallel Overhead



- $t$ : Serial time
- $n$  Number of processes
- $t_n$ : single chunk time with  $n$  chunks
- $t_p$ : Parallel time
- $t_s$ : Serial section time
- $t_o(n)$ : Parallel overhead time

# Preview - Distributed Compute Replication

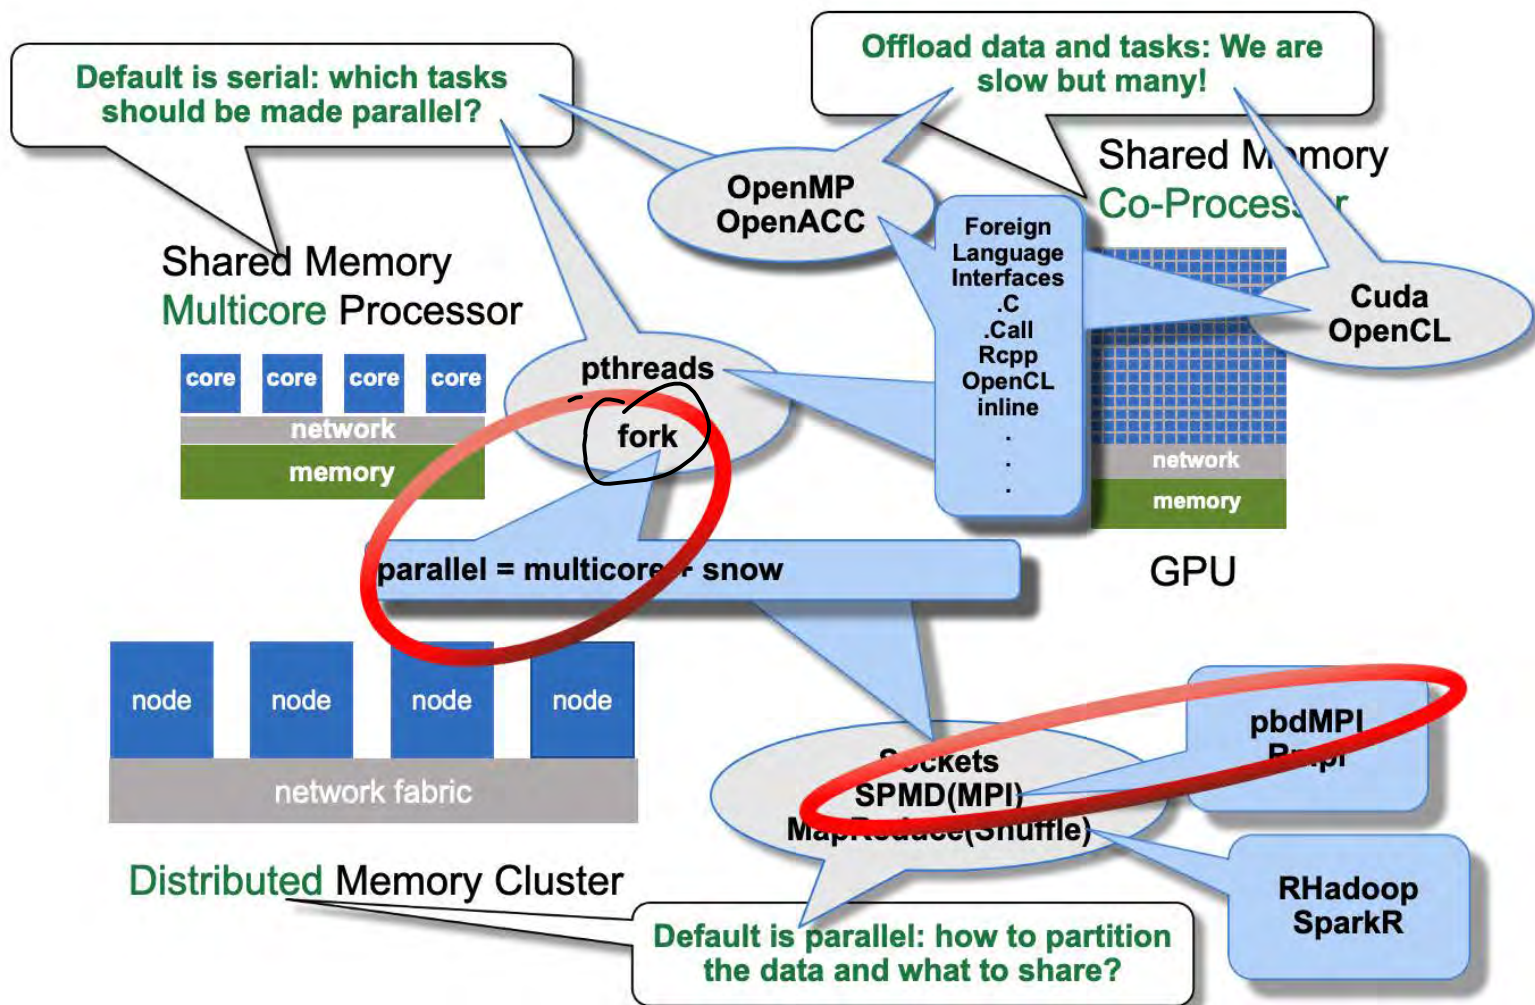


- $t$ : Serial time
- $n$  Number of processes
- $t_n$ : single chunk time with  $n$  chunks
- $t_p$ : Parallel time
- $t_s$ : Serial section time
- $t_o(n)$ : Parallel overhead time

**Replication can reduce communication overhead**



# R Interfaces to Low-Level Native Tools




# Unix fork


- Copy-on-write
- A memory efficient parallelism on shared memory devices
- **parallel** package `mcclapply` and friends
- Use for numerical sections only
  - Avoid GUI, I/O, and graphics sections
- Convenient for data (not modified)
- Convenient for functional languages like R
- Avoid or manage nested parallelism
  - OpenBLAS takes all cores by default
  - `data.table` automatically switches to single threaded mode upon fork

A deeper discussion of fork memory (if you have interest) on [YouTube](#) by Chris Kanich (UIC)

# Copy-on-write



Shared memory after forking Process B



After Process B writes to copy-on-write shared memory page

# Mapping Threads to Cores

## Theory and Reality

- Operating system manages core affinity
- Operating system tasks can compete
- Core switching occurs frequently
- **But it works rather well!**



## Drop-in replacements (almost) for `lapply`, `mapply`, and `Map`

```
mclapply(X, FUN, ..., mc.preschedule = TRUE, mc.set.seed =  
TRUE, mc.silent = FALSE, mc.cores = getOption("mc.cores", 2L),  
mc.cleanup = TRUE, mc.allow.recursive = TRUE, affinity.list =  
NULL)
```

```
mcmapply(FUN, ..., MoreArgs = NULL, SIMPLIFY = TRUE, USE.NAMES  
= TRUE, mc.preschedule = TRUE, mc.set.seed = TRUE, mc.silent =  
FALSE, mc.cores = getOption("mc.cores", 2L), mc.cleanup =  
TRUE, affinity.list = NULL)
```

```
mcMap(f, ...)
```



# Example Random forest Code

Letter recognition data ( 20 000  $\times$  17 )



```
[,1] lettr capital letter  
[,2] x.box horizontal position of box  
[,3] y.box vertical position of box  
[,4] width width of box  
[,5] high height of box  
[,6] onpix total number of on pixels  
[,7] x.bar mean x of on pixels in box  
[,8] y.bar mean y of on pixels in box  
[,9] x2bar mean x variance  
[,10] y2bar mean y variance  
[,11] xybar mean x y correlation  
[,12] x2ybr mean of x^2 y  
[,13] xy2br mean of x y^2  
[,14] x.ege mean edge count left to right  
[,15] xegvy correlation of x.ege with y  
[,16] y.ege mean edge count bottom to top  
[,17] yegvx correlation of y.ege with x
```

Figure 1: Letter Recognition data (image: [Frey and Slate, 1991](#)), description: **mlbench** package).

\*Parallel Statistical Computing with R: An Illustration on Two Architectures  
[arXiv:1709.01195](#)


# Random Forest Classification

Build many decision trees from random subsets of variables

Use their majority votes to classify

# Example Random Forest Classification Code

Letter recognition data ( 20 000  $\times$  17 )






# KPMS-IT4I-EX/code/rf\_serial.r

```
library(randomForest)
data(LetterRecognition, package = "mlbench")
set.seed(seed = 123)

n = nrow(LetterRecognition)
n_test = floor(0.2 * n)
i_test = sample.int(n, n_test)
train = LetterRecognition[-i_test, ]
test = LetterRecognition[i_test, ]

rf.all = randomForest(lettr ~ ., train, ntree = 500, norm.votes = FALSE)
pred = predict(rf.all, test)

correct = sum(pred == test$lettr)
cat("Proportion Correct:", correct/(n_test), "\n")
```



# KPMS-IT4I-EX/code/rf\_mc.r

```
library(parallel) —
library(randomForest)
data(LetterRecognition, package = "mlbench")
set.seed(seed = 123, "L'Ecuyer-CMRG")

n = nrow(LetterRecognition)
n_test = floor(0.2 * n)
i_test = sample.int(n, n_test)
train = LetterRecognition[-i_test, ]
test = LetterRecognition[i_test, ] ✓

nc = as.numeric(commandArgs(TRUE)[1])
ntree = lapply(splitIndices(500, nc), length)
rf = function(x) randomForest(lettr ~ ., train, ntree=x, norm.votes =
rf.out = mclapply(ntree, rf, mc.cores = nc)
rf.all = do.call(combine, rf.out)

crows = splitIndices(nrow(test), nc)
rfp = function(x) as.vector(predict(rf.all, test[x, ]))
cpred = mclapply(crows, rfp, mc.cores = nc)
pred = do.call(c, cpred)

correct <- sum(pred == test$lettr)
cat("Proportion Correct:", correct/(n_test), "\n")
```


parallel RNG

# Pseudo Random Number Generators (RNG)


- Guaranteed reproducibility
- Possibly overlapping streams

- Reproducibility for same number of streams
- Guaranteed independent streams

Random seeds with serial RNG



Parallel RNG



# KPMS-IT4I-EX/code/rf\_mc.r

```
library(parallel)
library(randomForest)
data(LetterRecognition, package = "mlbench")
set.seed(seed = 123, "L'Ecuyer-CMRG")

n = nrow(LetterRecognition)
n_test = floor(0.2 * n)
i_test = sample.int(n, n_test)
train = LetterRecognition[-i_test, ]
test = LetterRecognition[i_test, ]

nc = as.numeric(commandArgs(TRUE)[1])
ntree = lapply(splitIndices(500, nc), length)
rf = function(x) randomForest(lettr ~ ., train, ntree=x, norm.votes =
rf.out = mclapply(ntree, rf, mc.cores = nc)
rf.all = do.call(combine, rf.out)

crows = splitIndices(nrow(test), nc)
rfp = function(x) as.vector(predict(rf.all, test[x, ]))
cpred = mclapply(crows, rfp, mc.cores = nc)
pred = do.call(c, cpred)

correct <- sum(pred == test$lettr)
cat("Proportion Correct:", correct/(n_test), "\n")
```

*gets parameter from shell script*  
*rcp (125, 4)*  
*4 cores → 125 each*  
*returns a list*  
*overhead*  
*overhead*

# KPMS-IT4I- EX/code/rf\_karolina\_pbs.sh

```
#!/bin/bash
#PBS -N rf
#PBS -l select=1:ncpus=128,walltime=00:05:00
#PBS -q qexp
#PBS -e rf.e
#PBS -o rf.o

cd ~/KPMS-IT4I-EX/code
pwd

module load R
echo "loaded R"

time Rscript rf_serial.r
time Rscript rf_mc.r 1
time Rscript rf_mc.r 2
time Rscript rf_mc.r 4
```

# Demo ...