

# Adventures in Supercomputing with R

## Lecture 4: Shared Memory (continued)

George Ostrouchov

Oak Ridge National Laboratory and University of Tennessee

2022/02/28 (updated: 2022-03-09)

# Questions and Review

- RStudio - GitHub - It4I workflow?
- Are you using another workflow?
- IT4I experience?
- Measurement and terminology of parallel speedup
- Unix fork and its use by `mclapply()`
- Random Forest Classification

# Resampling ... and Crossvalidation

More shared memory parallelization, and a preview of parallel linear algebra ...

# Bootstrap, Bagging, Boosting, and Crossvalidation

- **Bootstrap**: a tool for assessing statistical accuracy
  - Resampling data with replacement and repeating estimation
  - Results in a sample of parameter estimates
- **Bagging** (bootstrap aggregation): a tool for reducing the variance of a prediction function
  - Simple models (low bias and high variance models) on resampled data
  - Consensus prediction (majority vote or average)
  - Generalized by **random forest** to sampling subsets of predictors
- **Boosting**: forward additive modeling, a greedy method of growing a model
  - Introduced via increasing weights on misclassified observations but shown\* to fit additive model framework
  - Sequential, so parallelization within a model
- **Crossvalidation**: model performance assessment
  - Estimates expected prediction error
  - Uses all data (no test set)

\*Hastie, Tibshirani, and Friedman (2009) The Elements of Statistical Learning, Second Edition, (2009). [link](#)

# Bootstrap

- Data:  $\mathbf{Z} = (z_1, z_2, \dots, z_N)$ , where  $z_i = (y_i, x_i)$
- Model: Let  $S(\mathbf{Z})$  be an estimated quantity from the data
- Sample with replacement  $B$  sets of size  $N$  from data
- Fit model to each of the resampled  $B$  sets  
 $\{S(\mathbf{Z}^{*1}), S(\mathbf{Z}^{*2}), \dots, S(\mathbf{Z}^{*B})\}$
- Use as sample from the sampling distribution of the estimator
- For example,  $\widehat{\text{Var}}[S(\mathbf{Z})] = \frac{1}{B-1} \sum_{b=1}^B [S(\mathbf{Z}^{*b}) - \bar{S}^*]^2$

Easy to parallelize over the  $B$  sets

Further opportunities may exist within the estimator  $S(\cdot)$

# Bagging (bootstrap aggregation)

- Let  $\widehat{S}(\mathbf{Z}) = \frac{1}{B} \sum_{b=1}^B S(\mathbf{Z}^{*b})$
- Majority vote if discrete
- Reduces variability of the estimate

Easy to parallelize over  $B$

Further opportunities may exist within the estimator  $S(\cdot)$

# Random Forest for Regression or Classification

1. For  $b = 1$  to  $B$ :
  - Draw a bootstrap sample  $\mathbf{Z}^*$  of size  $N$  from the training data.
  - Grow a random-forest tree  $T_b$  on  $\mathbf{Z}^*$ :
    - Recursively, for each terminal node, until  $n_{min}$  node size:
      - Select  $m$  variables at random from the  $p$  variables
      - Pick the best variable/split-point among the  $m$
      - Split the node into two daughter nodes
2. Output the ensemble of trees  $\{T_b\}_1^B$ .

To make a prediction at  $x$ :

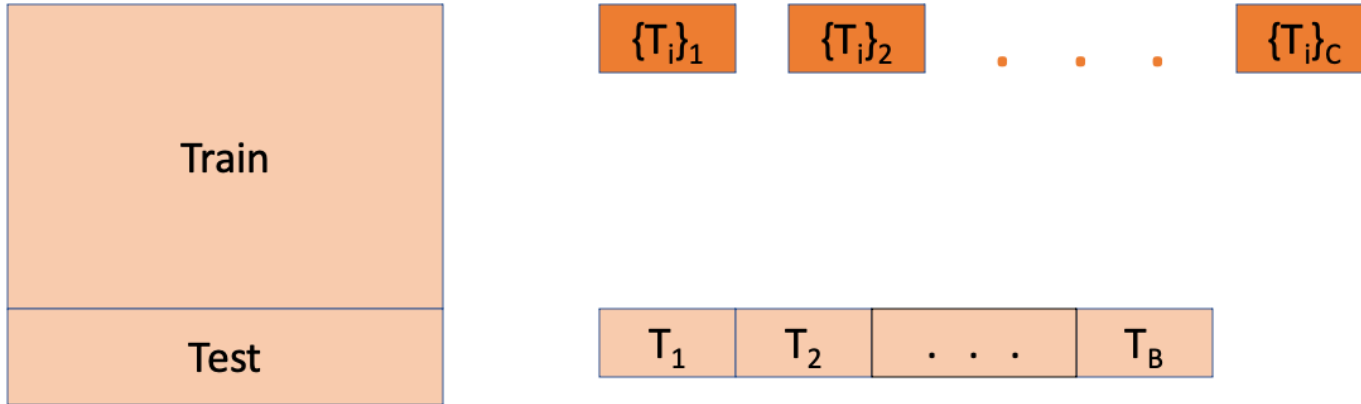
Regression:  $\hat{f}_{\text{rf}}^B(x) = \frac{1}{B} \sum_{b=1}^B T_b(x)$ .

Classification: Let  $\hat{C}_b(x)$  be the class prediction of the  $b$ th random-forest tree.

Then  $\hat{C}_{\text{rf}}^B(x) = \text{majority vote } \{\hat{C}_b(x)\}_1^B$ .

\*Algorithm 15.1 in Hastie, Tibshirani, and Friedman (2009). The Elements of Statistical Learning, Second Edition. [Link](#)

# Shared memory considerations



**Cores produce separate forests:**

**Combine forests for prediction or combine predictions?**



# Boosting

## Discrete AdaBoost \*

1. Initialize weights  $w_i = \frac{1}{N}, i = 1, 2, \dots, N$ .
2. For  $m = 1$  to  $M$ :
  - Fit a classifier  $G_m(x)$  to the training data using weights  $w_i$ .
  - Compute

$$err_m = \frac{\sum_{i=1}^N w_i I(y_i \neq G_m(x_i))}{\sum_{i=1}^N N w_i}$$

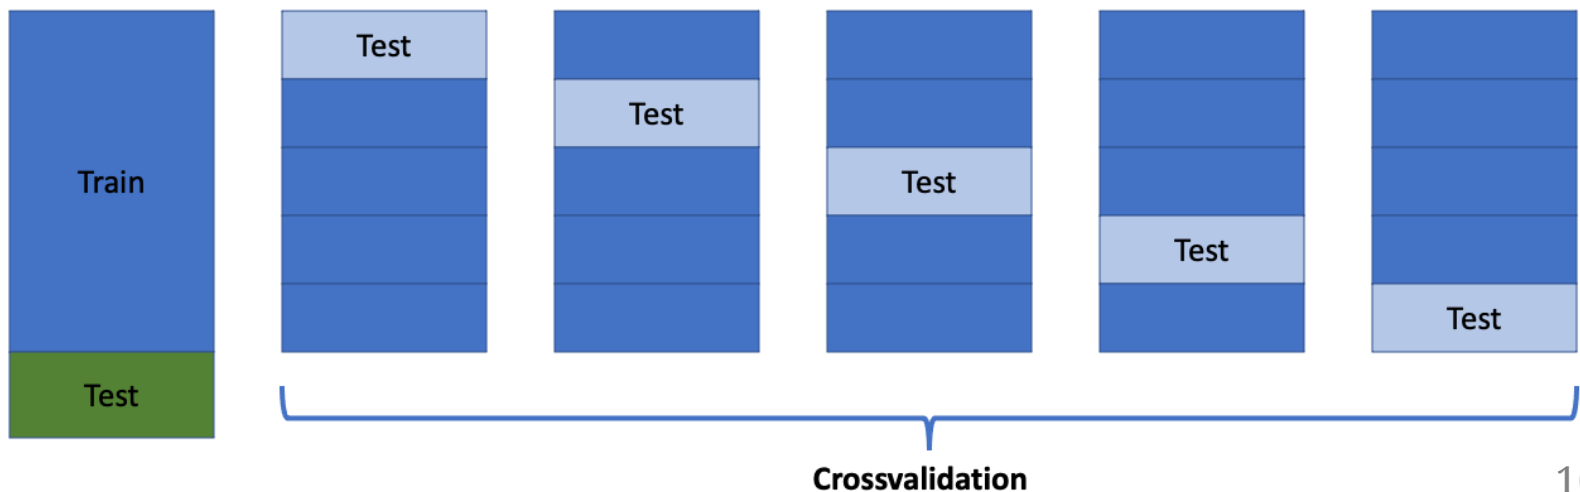
- Compute  $\alpha_m = \log((1 - err_m)/err_m)$ .
  - Set  $w_i \leftarrow w_i \cdot \exp[\alpha_m \cdot I(y_i \neq G_m(x_i))], i = 1, 2, \dots, N$ .
3. Output  $G(x) = \text{sign} \left[ \sum_{m=1}^M \alpha_m G_m(x) \right]$ .

## Sequential over M, parallelize within $G_m(\cdot)$

\* Algorithm 10.1 in Hastie, Tibshirani, and Friedman (2009)

# K-fold Crossvalidation

- Randomly divide data into  $k$  roughly equal folds
- Let  $\hat{f}^{-k}$  be the estimator when fold  $k$  is removed
- let  $k_i$  be the assigned fold of observation  $i$
- Let  $\hat{f}(x, \alpha)$  be an estimator with a tuning parameter  $\alpha$
- Then  $CV(f, \alpha) = \frac{1}{N} \sum_{i=1}^N L(y_i, \hat{f}^{-k_i}(x_i, \alpha))$  is the crossvalidation estimate of the prediction error



# KPMS-IT4I-EX/code/rf\_serial.r

```
library(randomForest)
data(LetterRecognition, package = "mlbench")
set.seed(seed = 123)

n = nrow(LetterRecognition)
n_test = floor(0.2 * n)
i_test = sample.int(n, n_test)
train = LetterRecognition[-i_test, ]
test = LetterRecognition[i_test, ]

rf.all = randomForest(lettr ~ ., train, ntree = 500, norm.votes = FALSE)
pred = predict(rf.all, test)

correct = sum(pred == test$lettr)
cat("Proportion Correct:", correct/(n_test), "\n")

plot(rf.all)
```

# KPMS-IT4I-EX/code/rf\_cv\_serial.r

```
set.seed(seed = 123)

ntree = 300
nfolds = 10
mtry_val = 1:(ncol(train) - 1)
folds = sample( rep_len(1:nfolds, nrow(train)), nrow(train) )
cv_df = data.frame(mtry = mtry_val, incorrect = rep(0, length(mtry_val)))
cv_pars = expand.grid(mtry = mtry_val, f = 1:nfolds)
fold_err = function(i, cv_pars, folds, train) {
  mtry = cv_pars[i, "mtry"]
  fold = (folds == cv_pars[i, "f"])
  rf.all = randomForest(lettr ~ ., train[!fold, ], ntree = ntree,
                        mtry = mtry, norm.votes = FALSE)
  pred = predict(rf.all, train[fold, ])
  sum(pred != train$lettr[fold])
}

cat("Running serial\n")
system.time({
  cv_err = lapply(1:nrow(cv_pars), fold_err, cv_pars, folds = folds, train = train)
  cv_err = tapply(unlist(cv_err), cv_pars[, "mtry"], sum)
})
png(paste0("rf_cv_mc0.png")); plot(mtry_val, cv_err/(n - n_test)); de
```

# KPMS-IT4I-EX/code/rf\_cv\_mc.r

```
set.seed(seed = 123, "L'Ecuyer-CMRG")

ntree = 200
nfolds = 10
mtry_val = 1:(ncol(train) - 1)
folds = sample( rep_len(1:nfolds, nrow(train)), nrow(train) )
cv_df = data.frame(mtry = mtry_val, incorrect = rep(0, length(mtry_val)))
cv_pars = expand.grid(mtry = mtry_val, f = 1:nfolds)
fold_err = function(i, cv_pars, folds, train) {
  mtry = cv_pars[i, "mtry"]
  fold = (folds == cv_pars[i, "f"])
  rf.all = randomForest(lettr ~ ., train[!fold, ], ntree = ntree,
                        mtry = mtry, norm.votes = FALSE)
  pred = predict(rf.all, train[fold, ])
  sum(pred != train$lettr[fold])
}

nc = as.numeric(commandArgs(TRUE)[1])
cat("Running with", nc, "cores\n")
system.time({
  cv_err = parallel::mclapply(1:nrow(cv_pars), fold_err, cv_pars, folds,
                             train = train, mc.cores = nc)
  err = tapply(unlist(cv_err), cv_pars[, "mtry"], sum)
})
```

# Demo ...

# Matrix Libraries ...

# R-LAPACK-BLAS



- BLAS: Basic Linear Algebra Subroutines - A matrix multiplication library
  - vector-vector (Level-1), matrix-vector (Level-2), matrix-matrix (Level-3)
- LAPACK: dense and banded matrix decompositions and more
  - $LU$   $LL^T$   $QR$   $UDV^T$   $VD^2V^T$   $\|\cdot\|_p$
- Implementations: OpenBLAS, Intel MKL, Nvidia nvBLAS, Apple vecLib, AMD BLIS, Arm Performance Libraries
- **FlexiBLAS**: A BLAS and LAPACK wrapper library with runtime exchangable backends



# FlexiBLAS

```
library(flexiblas)
```

```
# check whether FlexiBLAS is available
```

```
flexiblas_avail()
```

```
#> [1] TRUE
```

```
# get the current backend
```

```
flexiblas_current_backend()
```

```
#> [1] "OPENBLAS-OPENMP"
```

```
# list all available backends
```

```
flexiblas_list()
```

```
#> [1] "NETLIB"                "__FALLBACK__"        "BLIS-THREADS"        "OPENBLAS-OPENMP"
```

```
#> [5] "BLIS-SERIAL"           "ATLAS"               "OPENBLAS-SERIAL"     "OPENBLAS-OPENMP"
```

```
#> [9] "BLIS-OPENMP"
```

```
# get/set the number of threads
```

```
flexiblas_set_num_threads(12)
```

```
flexiblas_get_num_threads()
```

```
#> [1] 12
```

<https://github.com/Enchufa2/r-flexiblas>

<https://cran.r-project.org/package=flexiblas>