

Adventures in Supercomputing with R

Lecture 7: SVD and Friends, Begin Distributed

George Ostrouchov

Oak Ridge National Laboratory and University of Tennessee

2022/02/20 (updated: 2022-03-30)

Exercise 4 Review ...

Exercise 6 Review ...

Truncated SVD as Regression Basis Vectors

Images as rows (observations)

Continuing with the MNIST data set of digit images from last time, here we construct A with each image as a row.

Let A be the matrix of n_A images of a single digit, the pixel values of each image as a row (an observation). We have image samples of each of 10 digits and we put them into separate matrices, like A , and consider use them for classification of new images.

The SVD of $\mathbf{A} = \mathbf{U}\mathbf{D}\mathbf{V}^T$. If \mathbf{u}_i and \mathbf{v}_i are the columns of \mathbf{U} and \mathbf{V} , respectively, then

$$A = \sum_{i=1}^p d_i u_i v_i^T.$$

and image in row i is $\sum_{j=1}^p (d_j u_{ij}) v_j$. The v_j are basis vectors constructed from the training data.

Truncated SVD as Regression Basis Vectors

From matrix approximation, we know that this SVD can be truncated to some $k \ll p$ components and still represent the matrix (and so the images) well.

The v_i are basis functions constructed from data. It is a set of orthogonalized "images", which are the regressors and the $d_i u_{ij}$ are the regression coefficients on the training data.

We can now look at classification of a new image of a digit by regressing it onto each of the 10 digit bases and classifying it into the category that fits best.

The tuning parameter k can be optimized with crossvalidation.

New Image Regression onto the Basis Vectors

Suppose we performed SVDs on the ten matrices A_d , each being a training set of a digit $d = 0, 1, \dots, 9$, and produced a set of $p \times k$ matrices of k basis vectors \tilde{V}_d .

For a new image vector y , our digit d basis function model is $y = \tilde{V}_d \beta + \epsilon$. We drop the subscript d as we do the same for each digit.

Because the basis vectors are orthogonal, the normal equations $\tilde{V}^T \tilde{V} \hat{\beta} = \tilde{V}^T y$ are easily solved with $\hat{\beta} = \tilde{V}^T y$.

Note that because of the column truncation, we have $\tilde{V}^T \tilde{V} = I_k$ but $\tilde{V} \tilde{V}^T \neq I_p$. The L_2 loss is $\|y - \tilde{V} \tilde{V}^T y\|_2$.

For classification of the new image vector y , we choose the model with the smallest y regression model loss.

SVD and Eigenvalue Decompositions for Principal Components Analysis

In R, `princomp()` does an eigenvalue decomposition of the covariance (or correlation) matrix.

`prcomp()` computation proceeds via the singular value decomposition of the centered (and possibly scaled) data matrix.

$\text{Cov}(X)$ centers the data $X_c = X - n^{-1}11^T X$ and returns $X_c^T X$.

For $X = UDV^T$, note that

$$X^T X = VDU^T UDV^T = VD^2V^T$$

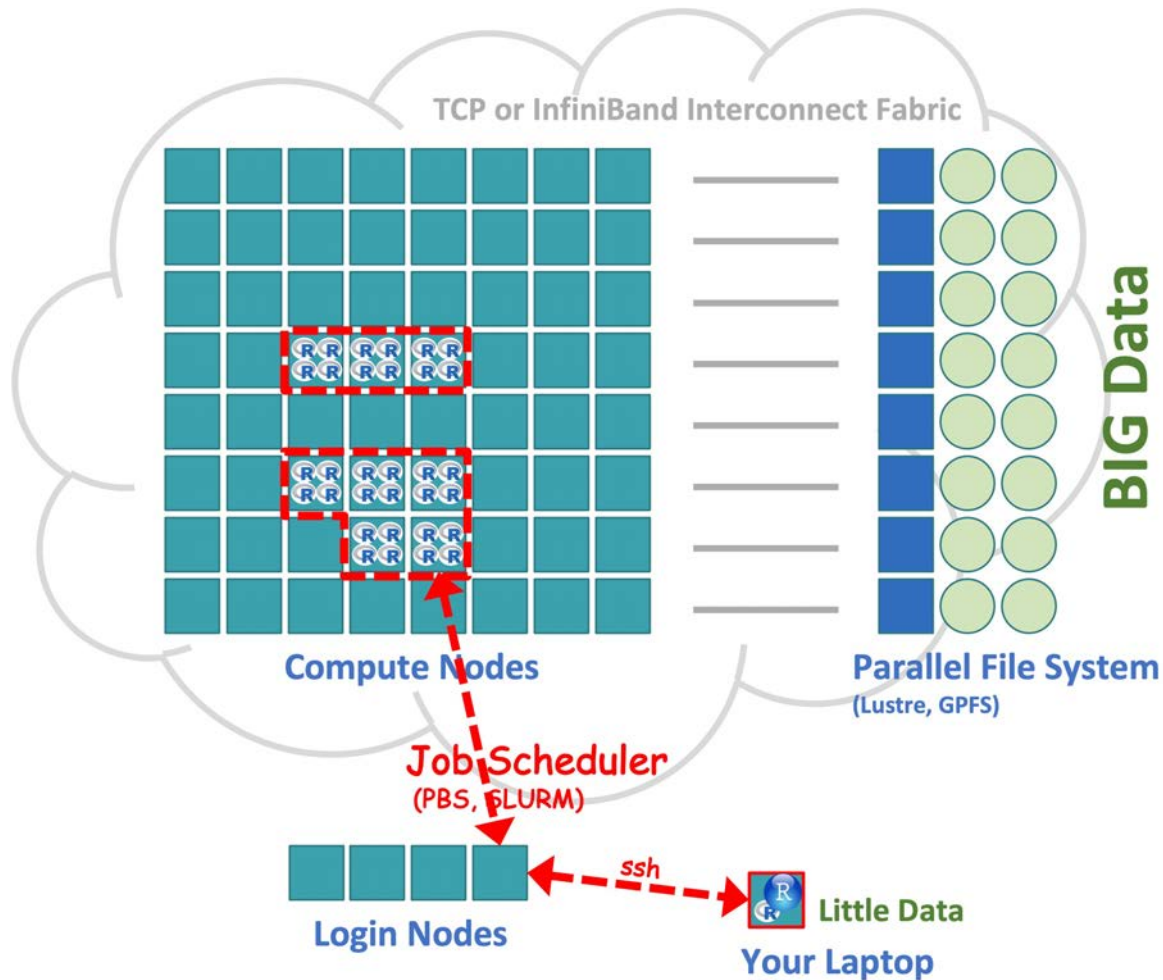
$$XX^T = UDV^T VDU^T = UD^2U^T$$

So that if $m \ll n$ or $m \gg n$ (a skinny X matrix), we can choose the smaller decomposition and recover the other $U = XV D^{-1}$ or $V = X^T U D^{-1}$ with matrix multiplication.

This results in easy fast distributed algorithms for PCA of skinny matrices as we will see in a future lecture.

Questions? ... Demo ...

Running Distributed on a Cluster



Parallel Computing Models

Shared memory parallel computing

- Processors have access to all memory
 - Locking mechanism
- Kinds: unix fork, pthreads, OpenMP, OpenACC
- Libraries: OpenBLAS, MKL, FlexiBLAS, PLASMA, MAGMA, etc.

Distributed memory parallel computing

- Processors have only local memory
 - Communication mechanism
- Kinds: MPI, MapReduce, DataFlow
- Libraries: OpenMPI, ScaLAPACK, PETSci, Trillinos, etc.

Distributed Programming Models


- Manager-workers
 - Common in simple cases (e.g. `mclapply()`)
- Single program, multiple data (SPMD)
 - Most common on supercomputer clusters
 - Most scalable because based on collaboration rather than control
 - Concept behind MPI design
- MapReduce
 - Common in data processing but falling out of favor because
 - Slow because includes all-to-all shuffle communication
- Dataflow
 - Dependency graph directed, still evolving
 - Sometimes combined with MPI for distributed control

pbdR-ScaLAPACK-PBLAS-BLACS-MPI



- MPI: Message Passing Interface - *de facto* standard for distributed communication in supercomputing
 - Used for data mostly via collective communication - high level
 - pbdMPI, kazaam, and cop R packages
- ScaLAPACK: Scalable LAPACK - Distributed version of LAPACK (uses PBLAS/BLAS but not LAPACK)
 - 2d Block-Cyclic data layout - mostly automated in pbdDMAT package
 - BLACS: Communication collectives for distributed matrix computation
 - PBLAS: BLAS - distributed BLAS (uses shared memory BLAS within blocks)
 - pbdDMAT and pbdML R packages - most matrix operations identical to serial through overloading operators and `ddmatrix` class

R Interfaces to Low-Level Native Tools



Single Program Multiple Data (SPMD)

Hello world!

```
suppressMessages(library(pbdMPI))  
  
msg = paste("Hello World! rank", comm.rank(), "out of", comm.size())  
cat(msg, "\n")  
  
finalize()
```

One code and a parallel mindset

A generalization of a serial code

Many rank-aware operations are automated

No manager, it is all cooperation

Explicit point-to-point communications are an advanced topic

pbdMPI:

reduce(X)

allreduce(X)

Computing Crossproducts on 8 Processors of the First (1986) Intel IPSC/1 Hypercube

$$A = X^T X = \sum_i X_i^T X_i, \text{ where } X = \begin{bmatrix} X_1 \\ \vdots \\ X_8 \end{bmatrix}$$




FIG. 4. Computation of $A = X^T X$ on an 8-processor hypercube, with final result on processor 0.




FIG. 6. Computation of $A = X^T X$ on an 8-processor hypercube, with final result on all processors.